

# AKS – Best Practices

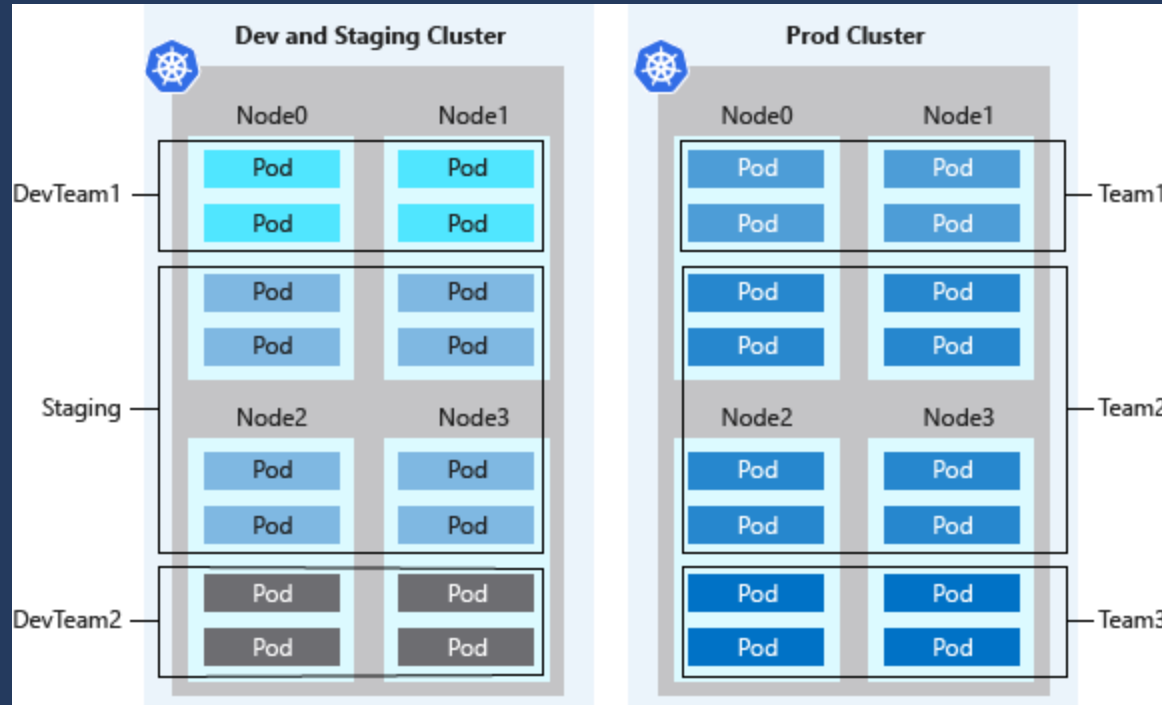


- Multi-Tenancy
  - Cluster isolation
  - Scheduler Features
  - Authentication & Authorization
- Security
  - Cluster upgrades
  - Container image management
  - Pod security
- Network and Storage
  - Network Connectivity
  - Storage and Backups
- BCP DR

# Cluster Isolation

- Logical Isolation

- A single AKS cluster can be used for multiple workloads, teams, or environments
- higher pod density than
- **Namespace** creates a logical isolation boundary

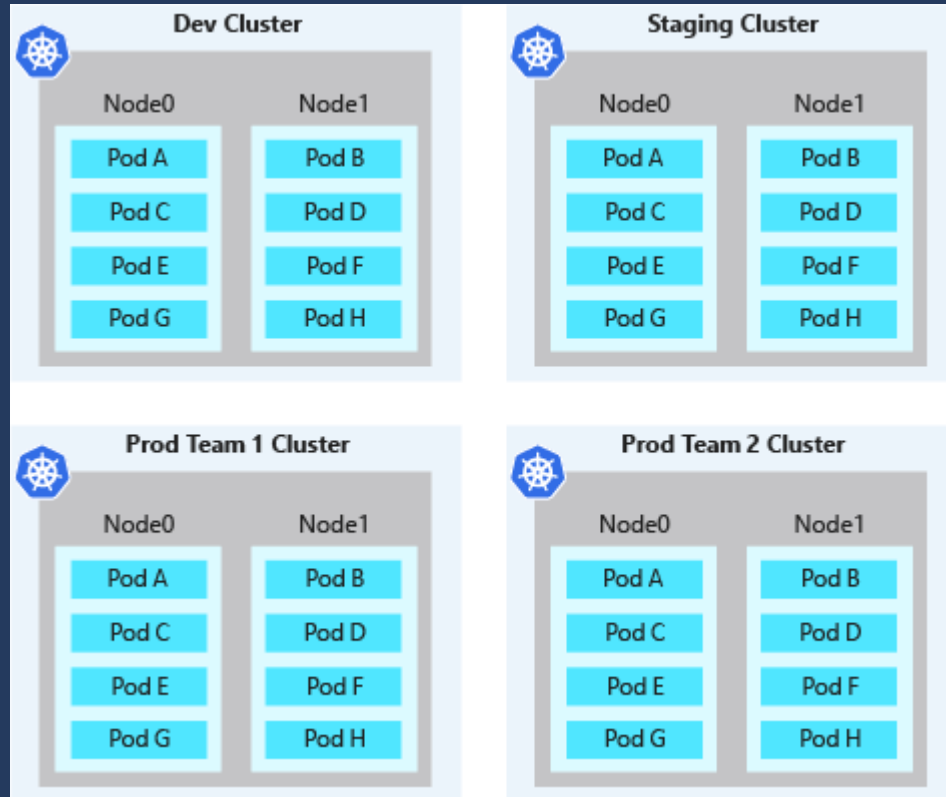


**Cluster Isolation**

**Logical Isolation**

Not completely safe for hostile multi tenant usage

- Teams or workloads are assigned their own AKS cluster
- Adds management and financial overhead



**Cluster Isolation**

**Physical Isolation**

- Minimize the use of physical isolation for each separate team or application deployment

# Scheduler Features

# Default Requests and Limits

- In a large system, it may not be practical to define resource requests/limits for all the existing containers.
- The **LimitRange** object allows you to specify a default set of requests (***defaultRequests***) and limits (***defaults***) for any Container created in a *Namespace*.
- When requests/limits are not specified by the Container, the defaults in the **LimitRange** will be used.
- When requests/limits are specified, those values override the defaults in the **LimitRange**.
- The ***min*** (requests) and ***max*** (limits) can also be specified to ensure values specified for a Container stay within set bounds.

```
apiVersion: v1
kind: LimitRange
metadata:
  name: mem-limit-range
spec:
  limits:
    # Default limit if not specified
    - default:
        cpu: 200m
        memory: 256Mi
      # Default request if not specified
      defaultRequest:
        cpu: 100m
        memory: 128Mi
      # Max limit if specified
      max:
        cpu: 1
        memory: 1Gi
      # Min request if specified
      min:
        cpu: 50m
        memory: 100Mi
    type: Container
```

# Enforce resource quotas

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: dev-app-team
spec:
  hard:
    cpu: "10"
    memory: 20Gi
    pods: "10"
```

```
kubectl apply -f dev-app-team-quotas.yaml --
namespace dev-apps
```

Plan and apply resource quotas at the namespace level. If pods don't define resource requests and limits, reject the deployment. Monitor resource usage and adjust quotas as needed.

---

- **Compute resources**, such as CPU and memory, or GPUs.
- **Storage resources**, including the total number of volumes or amount of disk space for a given storage class.
- **Object count**, such as maximum number of secrets, services, or jobs can be created.



# Namespace Resource Quota

- A **ResourceQuota** object, defines constraints that limit aggregate resource (total CPU and memory) consumption per namespace.
- It's **highly recommended** to create a **LimitRange** resource in the same namespace as the **ResourceQuota**.
- Limits should be large enough to accommodate upgrades.

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: mem-cpu-rq
spec:
  hard:
    requests.cpu: "1200m"
    limits.cpu: "2400m"
    requests.memory: 1.5Gi
    limits.memory: 3Gi
```

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: resources-rq
spec:
  hard:
    pods: "10"
    configmaps: "10"
    secrets: "10"
    persistentvolumeclaims: "4"
    services: "10"
    services.loadbalancers: "2"
```

# Node Selectors

- Pods can use Node labels to specify which Nodes to be scheduled on.
- Labels must be on the Nodes prior to deploying Pods selecting the Nodes.
- If selected Node label is not found on any node, Pods will not get scheduled.

```
apiVersion: v1
kind: Node
metadata:
  annotations:
    node.alpha.kubernetes.io/ttl: "0"
    volumes.kubernetes.io/controller-managed-
  labels:
    kubernetes.io/os: linux
    size: large
    topology.kubernetes.io/region: us-east-1
    topology.kubernetes.io/zone: us-east-1a
```



```
spec:
  nodeSelector:
    kubernetes.io/os: linux
  containers:
    - name: nginx
      image: k8slab/nginx:1.0
      ports:
        - containerPort: 80
          protocol: TCP
```

# Node Affinity

- The affinity/anti-affinity language is more expressive. The language offers more matching rules besides exact matches created with a logical AND operation
- Can indicate that the rule is soft "preference" rather than a hard requirement, so if the scheduler can't satisfy it, the Pod will still be scheduled.
- Use ***weight*** to set preference order.

```
spec:
  affinity:
    nodeAffinity:
      requiredDuringSchedulingIgnoredDuringExecution:
        nodeSelectorTerms:
          - matchExpressions:
              - key: kubernetes.io/e2e-az-name
                operator: In
                values:
                  - e2e-az1
                  - e2e-az2
      preferredDuringSchedulingIgnoredDuringExecution:
        - weight: 1
          preference:
            matchExpressions:
              - key: another-node-label-key
                operator: In
                values:
                  - another-node-label-value
    containers:
      - name: with-node-affinity
        image: k8s.gcr.io/pause:2.0
```

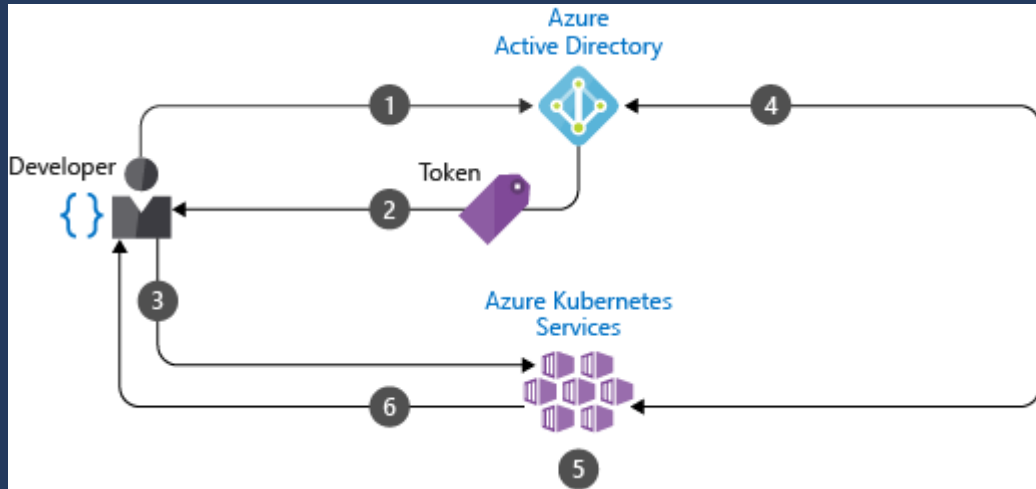
# Taints and Tolerations

- Since a cluster can be a collection of heterogeneous nodes, sometimes it makes sense to set a policy restricting which Pods are scheduled on which Nodes.
- Nodes can be marked as "tainted" with particular attributes
- Pods can be designated as being able to "tolerate" certain taints
- Depending on the taint-effect, the scheduler will decide if a Pod can be scheduled on a tainted Node.
- Node taint-effects can be:
  - **NoSchedule** – New Pods will NOT be scheduled on the tainted node unless they can tolerate the taint
  - **PreferNoSchedule** – Pods CAN be scheduled ONLY IF they won't fit on any other node.
  - **NoExecute** – New Pods will NOT be scheduled on the tainted node and existing Pods without toleration will be EVICTED from a tainted node.

# Authentication and Authorization



## Use Azure Active Directory (Azure AD)



# Authentication and Authorization

## Use Kubernetes role-based access control (Kubernetes RBAC)

## Use Azure RBAC

# Kubernetes Role-based Access Control (RBAC)

**RBAC** is a method of regulating access to resources based on the roles of individual users or service accounts within your organization.

RBAC Components:

- **Service Account** - Provides an identity for processes that run in a Pod.
- **Role** – A list of rights (permissions) to specific resource types within a namespace.
- **RoleBinding** – Defines the *binding* of a user/service account to a Role within a namespace
- **ClusterRole** – A cluster-wide resource listing permissions for specific namespaces, all namespaces or cluster-scoped resources.
- **ClusterRoleBinding** – Defines the binding of user/service account to a **ClusterRole** throughout the cluster

# Kubernetes Role-based Access Control (RBAC)

```
kind: Role
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: finance-app-full-access-role
  namespace: finance-app
rules:
- apiGroups: [""]
  resources: ["*"]
  verbs: ["*"]
```

```
kind: RoleBinding
apiVersion: rbac.authorization.k8s.io/v1
metadata:
  name: finance-app-full-access-role-binding
  namespace: finance-app
subjects:
- kind: User
  name: developer1@contoso.com
  apiGroup: rbac.authorization.k8s.io
roleRef:
  kind: Role
  name: finance-app-full-access-role
  apiGroup: rbac.authorization.k8s.io
```

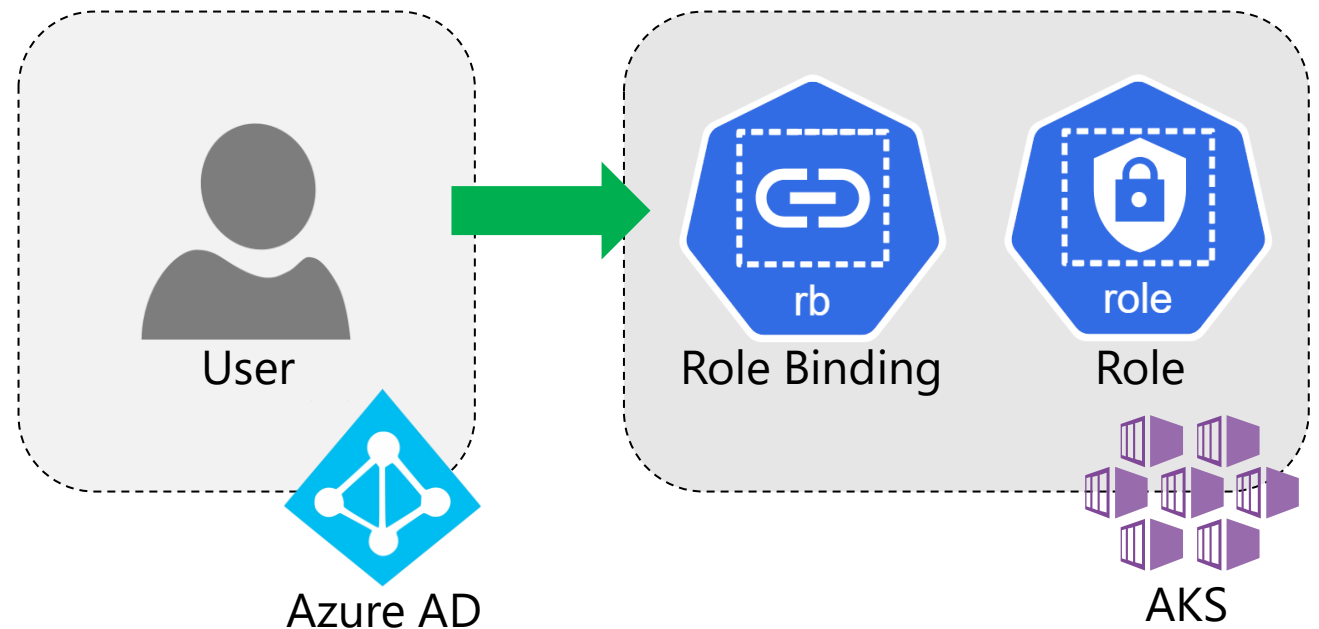
Two levels of access needed to fully operate an AKS cluster:

- Access the AKS resource on your Azure subscription.
  - Control scaling or upgrading your cluster using the AKS APIs
  - Pull your kubeconfig.
- Access to the Kubernetes API.
  - Kubernetes RBAC (traditionally) or
  - By integrating Azure RBAC with AKS for kubernetes authorization



# Kubernetes RBAC with Azure AD Integration

- Kubernetes RBAC is enabled by default in AKS. Disabling RBAC is **not recommended**.
- To enable Azure AD integration, use the `--enable-aad` option when creating or updating an AKS cluster.
- Kubernetes RBAC works with Azure AD as follows:
  1. A user is authenticated in Azure
  2. The user's token and email is passed to the cluster.
  3. A **RoleBinding** resource binds the user to a **Role**, which define which actions are allowed on which Kubernetes resources.



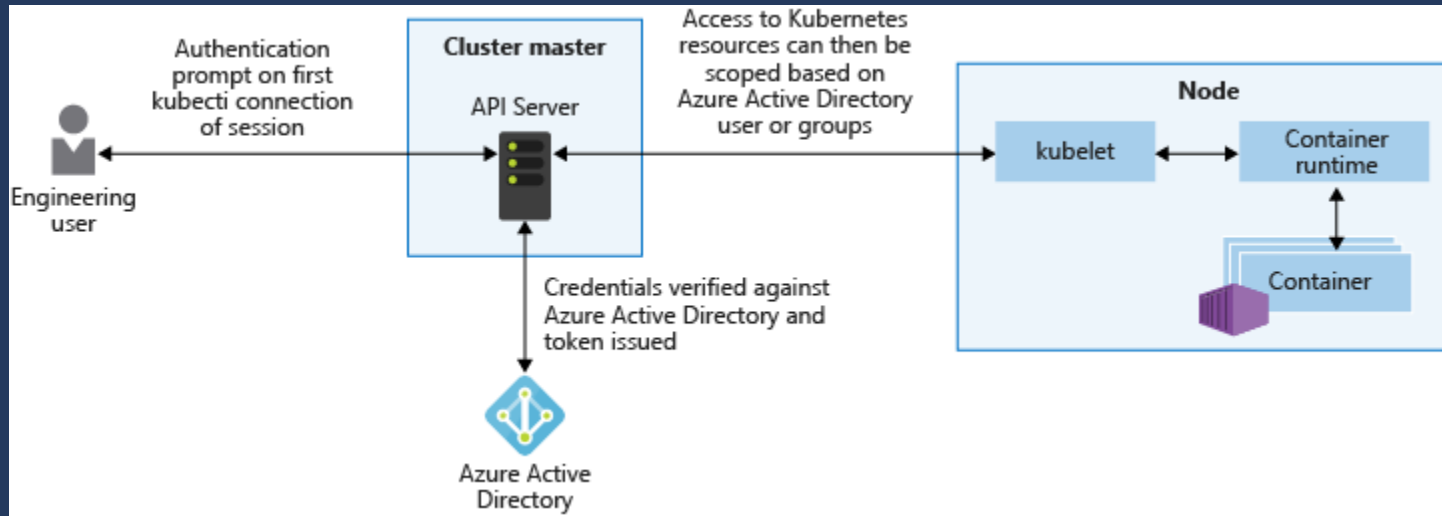
# Cluster Security & Upgrades

## Defender for Containers

Assess cluster configurations and provide security recommendations, run vulnerability scans, and provide real-time protection and alerting for Kubernetes nodes and clusters.

## Secure access to cluster and nodes

Integrate with K8s RBAC with Azure AD



## Threat protection

## Permissions to containers

Use minimum privileges , avoid root access or privileged escalations

App Armor

Seccomp

## Upgrade Kubernetes versions

Regularly upgrade the Kubernetes version

```
az aks get-upgrades --resource-group myResourceGroup --  
name myAKSCluster --output table
```

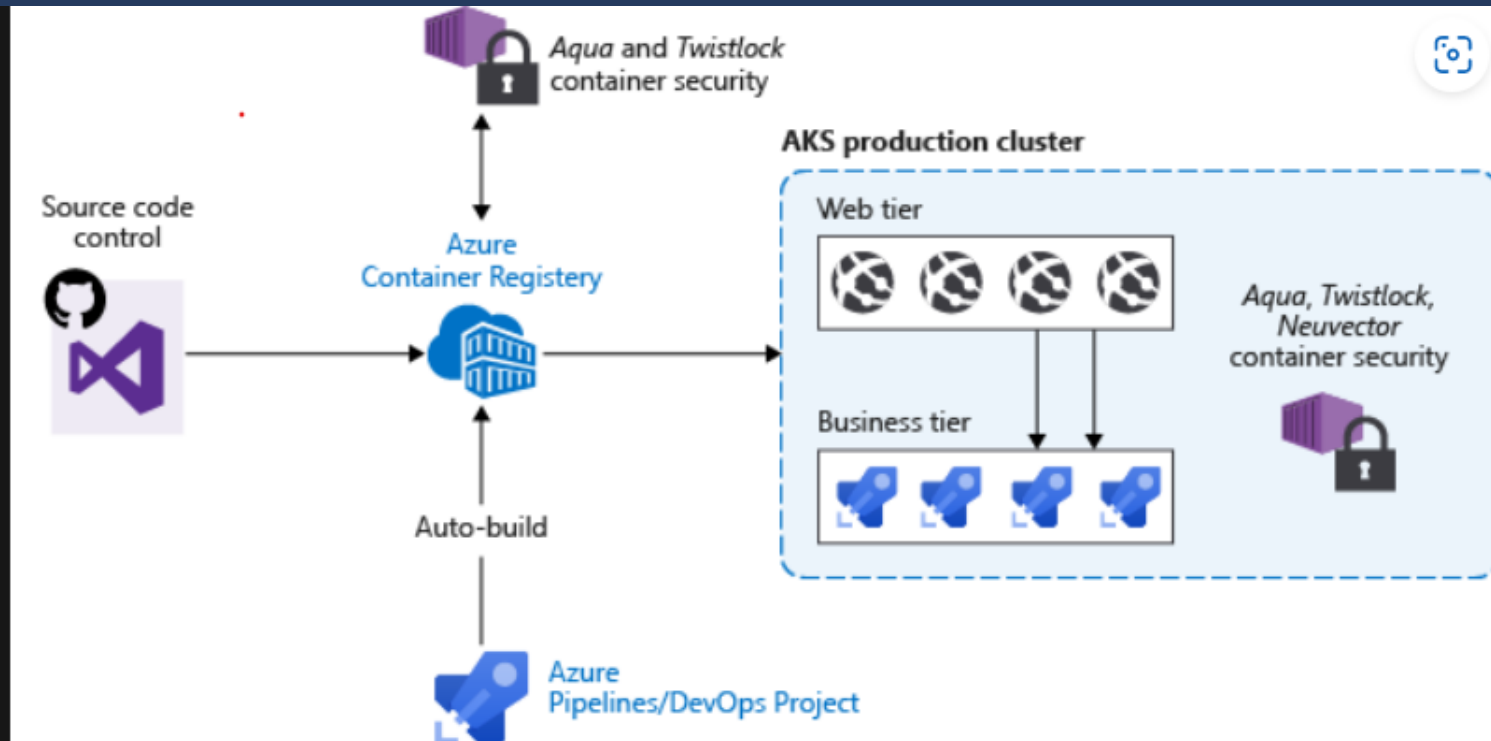
```
az aks upgrade --resource-group myResourceGroup --name  
myAKSCluster --kubernetes-version KUBERNETES_VERSION
```

## Secure Container Access

# Container Image management

- Scan your container images for vulnerabilities. Only deploy validated images.
- Include in your deployment workflow a process to scan container images using tools such as Twistlock or Aqua.
- 

## Build and Runtime Security



- Automatically build new images on base image update
  - Use automation to build new images when the base image is updated.
  - ACR Tasks can also automatically update container images when the base image is updated

#### Dockerfile-app

```
ARG REGISTRY_NAME
FROM ${REGISTRY_NAME}/baseimages/node:15-alpine

COPY . /src
RUN cd /src && npm install
EXPOSE 80
CMD ["node", "/src/server.js"]
```

#### Dockerfile-base

```
FROM node:15-alpine
ENV NODE_VERSION 15.2.1
```

## Container Image management

- Build base image

```
az acr build --registry $ACR_NAME --image baseimages/node:15-alpine --file Dockerfile-base .
```

- Create ACR task

```
az acr task create \  
  --registry $ACR_NAME \  
  --name baseexample1 \  
  --image helloworld:{{.Run.ID}} \  
  --arg REGISTRY_NAME=$ACR_NAME.azurecr.io \  
  --context https://github.com/$GIT_USER/acr-build-helloworld-node.git#main \  
  --file Dockerfile-app \  
  --git-access-token $GIT_PAT
```

- Manual Trigger

```
az acr task run --registry $ACR_NAME --name baseexample1
```

**Update node version :** ENV NODE\_VERSION 15.2.1a

```
az acr build --registry $ACR_NAME --image baseimages/node:15-alpine --file Dockerfile-base .
```

```
az acr task list-runs --registry $ACR_NAME --output table
```

Output							Copy
Run ID	TASK	PLATFORM	STATUS	TRIGGER	STARTED	DURATION	
ca11	baseexample1	linux	Succeeded	Image Update	2020-11-20T23:38:24Z	00:00:34	
ca10	taskhelloworld	linux	Succeeded	Image Update	2020-11-20T23:38:24Z	00:00:24	
cay		linux	Succeeded	Manual	2020-11-20T23:38:08Z	00:00:22	

- Now set the scheduled trigger (`--schedule "0 21 * * *" \` in acr task create command)

# Container Image management

Builds the base image, pushes it to container registry if the build is successful.



## Public Registries



Docker Hub



MCR



RedHat



gcr



GitHub

Internet



Private Registry

acr, ecr, gcr, dtr,  
harbor, jfrog, ...

registry.acme-rockets.io/

import-staging/

dotnet

node

alpine

nginx

base-artifacts/

dotnet

node

alpine

nginx

dev/

team-tasmania

team-gossamer

prod/

marketing

warranty

## Acme-Rockets

azure, aws, google, ibm, on-prem, ...

PIPELINE



TESTING

dotnet

node

alpine

nginx

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

node

alpine

nginx

team-tasmania

team-gossamer

marketing

warranty

PIPELINE

TESTING

dotnet

## Stable Tags :

*Tags that is reused, for example, major or minor version  
mycontainerimage:1.0.*

Use stable tags to maintain **base images** for your container builds

These tags continue to receive updates.

**Unique tags :** *A different tag for each image pushed to a registry, mycontainerimage:abc123.*

Use unique tags for **deployments**, especially in an environment that could scale on multiple nodes.

*Some patterns*

- *date-time*
- *git commit*
- *manifest digest*
- *build id*

## Lock deployed image tags

lock any deployed image tag, by setting its `write-enabled` attribute to `false`.

# Container Image Tagging

Multi-tenancy and cluster isolation

Logical Vs Physical

Scheduling features

Resource Quotas

Taint & Toleration

Node affinity

Networking

Azure CNI , Kubenet, Ingress traffic

Network Policies

Storage

Dynamic provisioning

Secure and backup - Velero

**For Operators**

# Pod Security

•**allowPrivilegeEscalation** Design your applications so this setting is always set to *false*.

•**Linux capabilities** let the pod access underlying node processes.

•**SELinux labels** is a Linux kernel security module that lets you define access policies for services, processes, and filesystem access.

- Pod runs as user ID *1000* and part of group ID *2000*
  - Can't escalate privileges to use root
  - Allows Linux capabilities to access network interfaces and the host's real-time (hardware) clock
- 

apiVersion: v1

kind: Pod

metadata:

name: security-context-demo

spec:

**securityContext:**

**fsGroup: 2000**

containers:

- name: security-context-demo

image: mcr.microsoft.com/oss/nginx/nginx:1.15.5-alpine

**securityContext:**

**runAsUser: 1000**

**allowPrivilegeEscalation: false**

**capabilities:**

**add: ["NET\_ADMIN", "SYS\_TIME"]**

Resource management

VS Code extension for Kubernetes

Debug with bridge to kubernetes

Pod Security

**For Developers**

# Networking and Storage

Choose the right network model  
(kubenet,Azure CNI)

Distribute ingress traffic – Ingress controllers

Secure traffic with a web application firewall

Control traffic flow with network policies

Securely connect to nodes through a bastion  
host

## **Networking**

# Network Policy

```
kind: NetworkPolicy
apiVersion: networking.k8s.io/v1
metadata:
  name: backend-policy
spec:
  podSelector:
    matchLabels:
      app: backend
  ingress:
    - from:
      - podSelector:
          matchLabels:
            app: frontend
```

Use network policies to allow or deny traffic to pods. By default, all traffic is allowed between pods within a cluster. For improved security, define rules that limit pod communication.

---

- You create a network policy as a Kubernetes resource using a YAML manifest.
- Policies are applied to defined pods, with ingress or egress rules defining traffic flow.
- As pods are dynamically created in an AKS cluster, required network policies can be automatically applied.

Reference:

[ahmetb/kubernetes-network-policy-recipes: Example recipes for Kubernetes Network Policies that you can just copy paste \(github.com\)](https://github.com/ahmetb/kubernetes-network-policy-recipes)



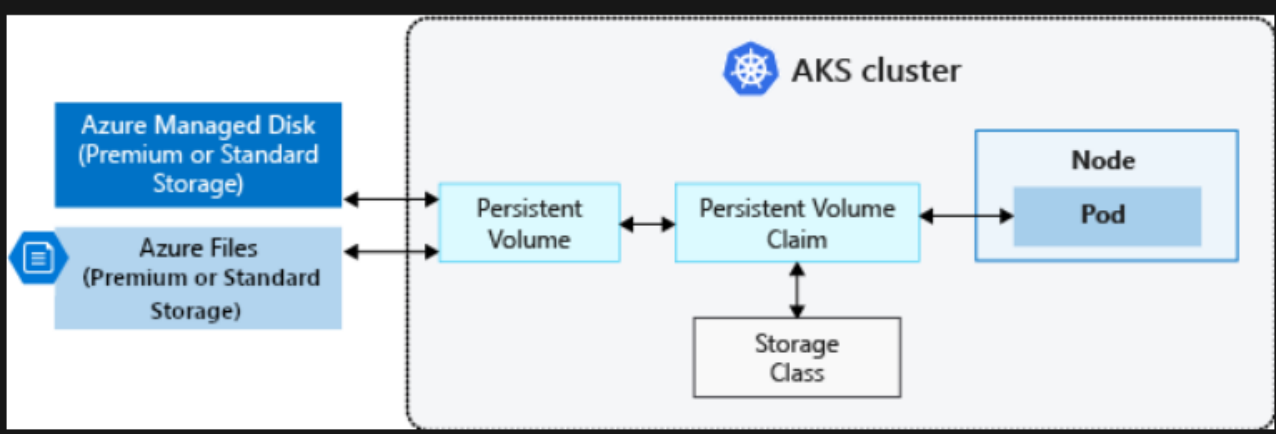
# Choose the right storage type

Use case	Volume plugin	Read/write once	Read-only many	Read/write many	Windows Server container support
Shared configuration	Azure Files	Yes	Yes	Yes	Yes
Structured app data	Azure Disks	Yes	No	No	Yes
Unstructured data, file system operations	<a href="#">BlobFuse</a>	Yes	Yes	Yes	No

## Storage

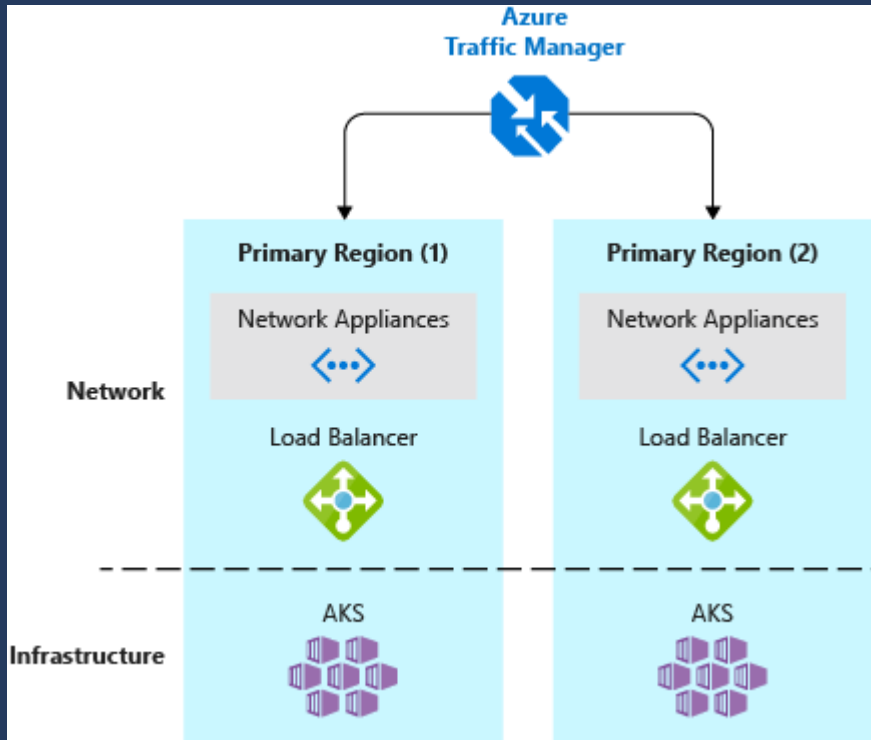
Backup data with Velero or Azure backup

# Dynamic volume provisioning



# Business Continuity and Disaster Recovery

- Plan for multi region deployment
- Use Azure Traffic manager to route request

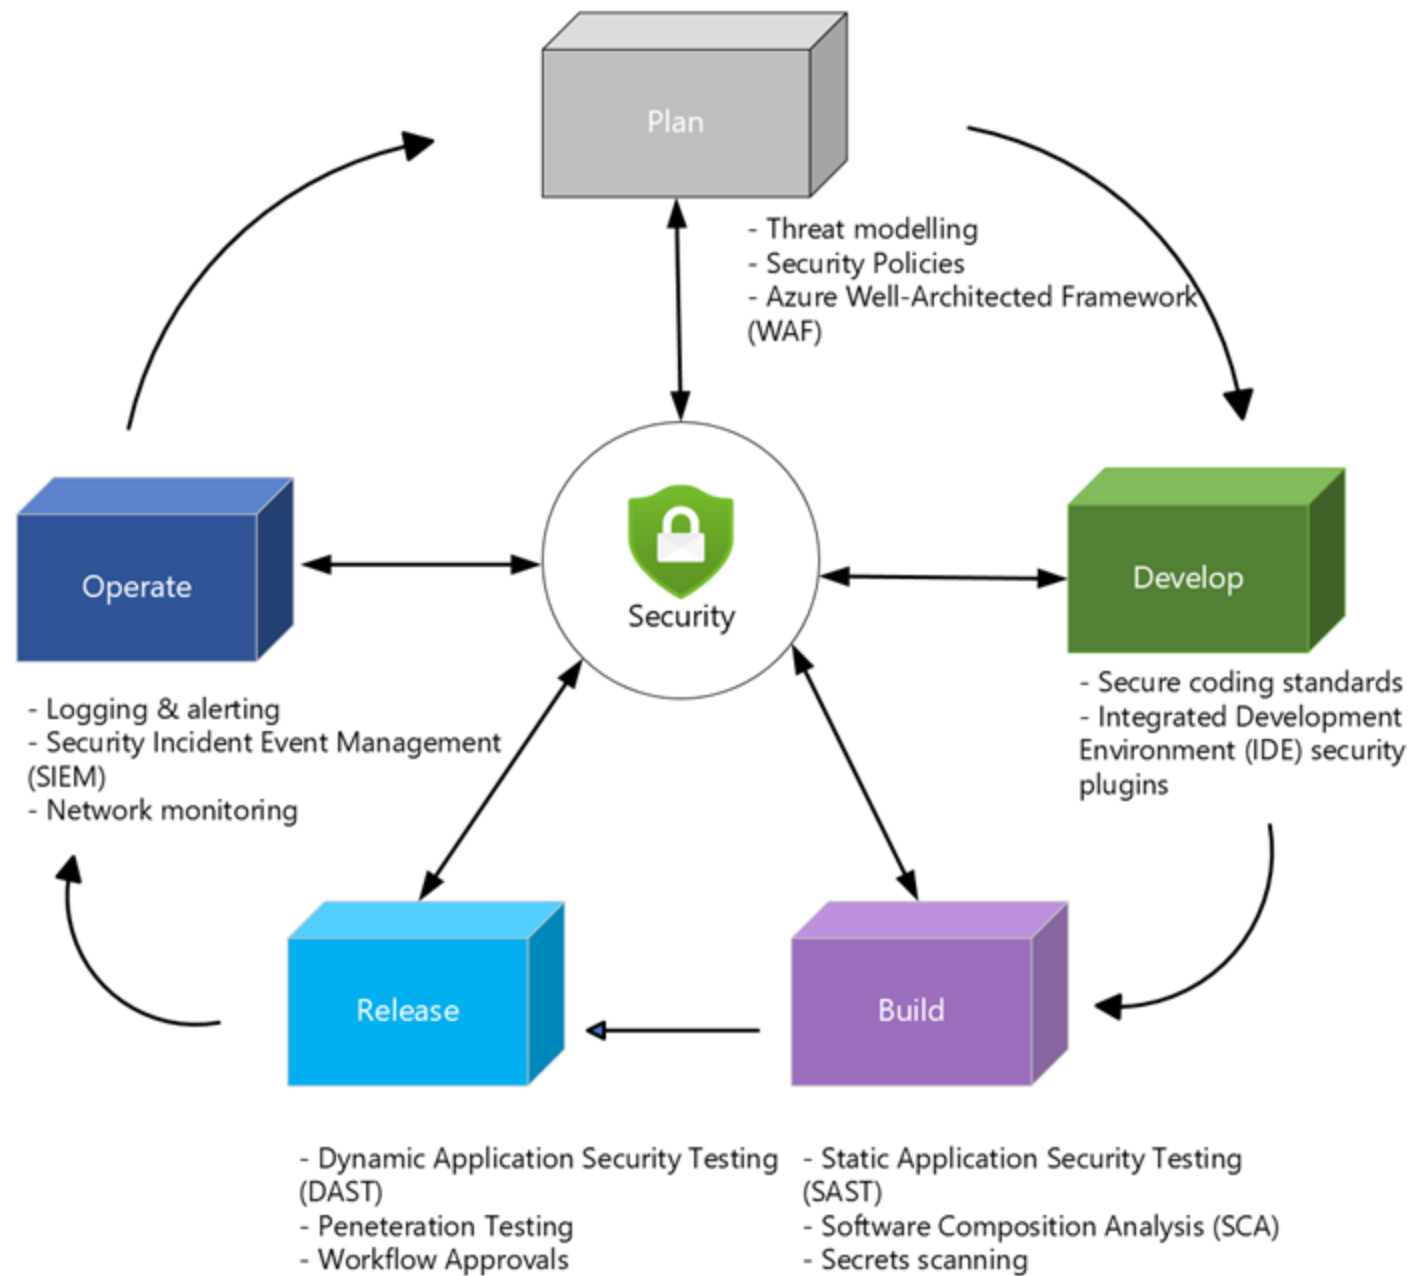


**BCP/DR**

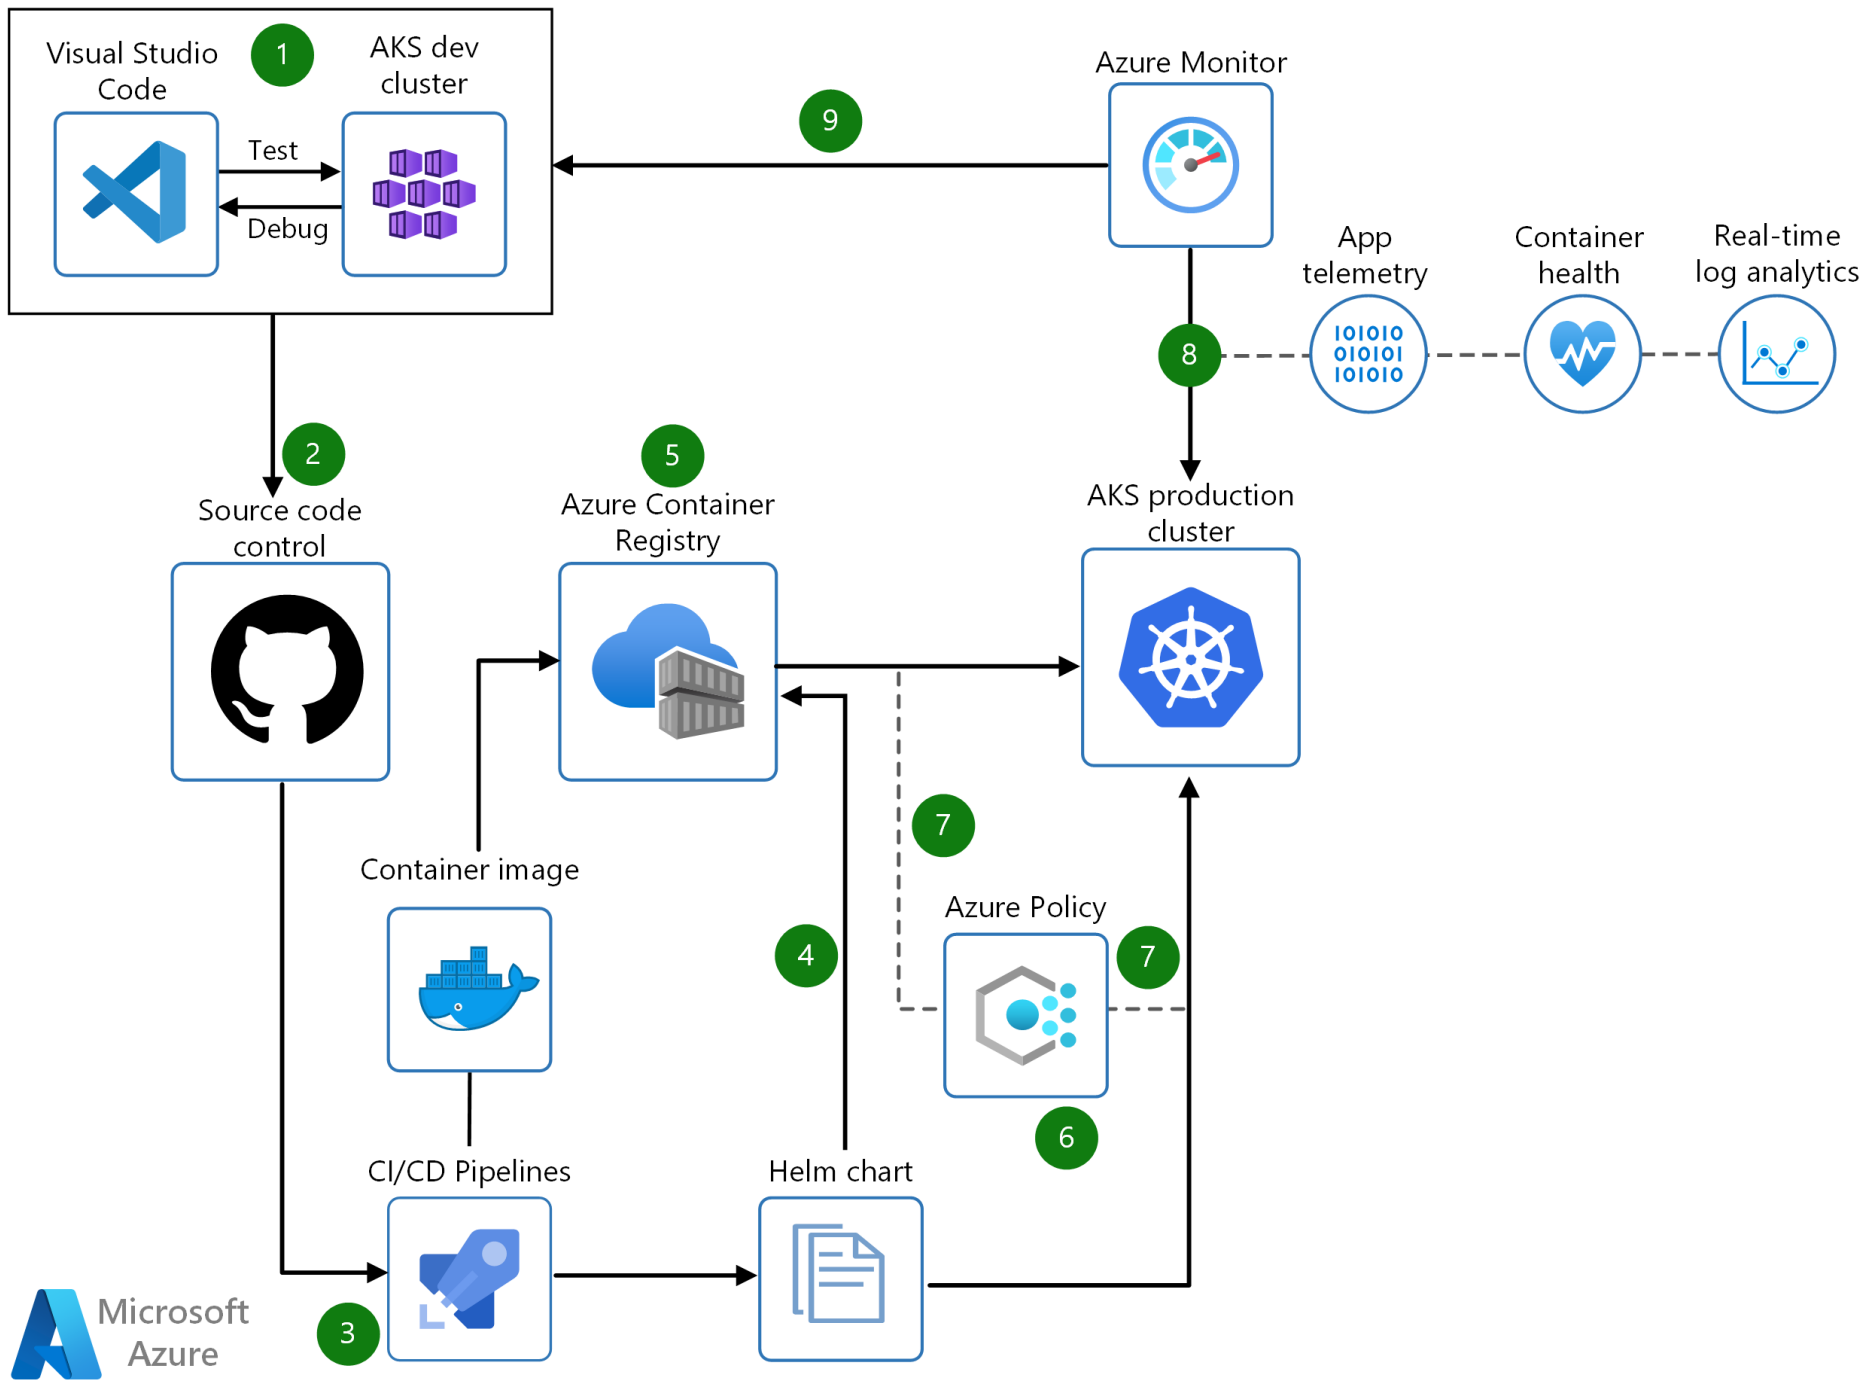
- Enable Geo-replication for container images

# DevSecOps in AKS





## Inner loop




 [Configure Azure CNI networking in Azure Kubernetes Service \(AKS\) - Azure Kubernetes Service | Microsoft Docs](#)

[Baseline architecture for an AKS cluster - Azure Architecture Center | Microsoft Docs](#)

[11 Ways \(Not\) to Get Hacked | Kubernetes](#)

[The Azure Kubernetes Service Checklist - ✨ Be ready for production ✨ \(the-aks-checklist.com\)](#)







Thank you